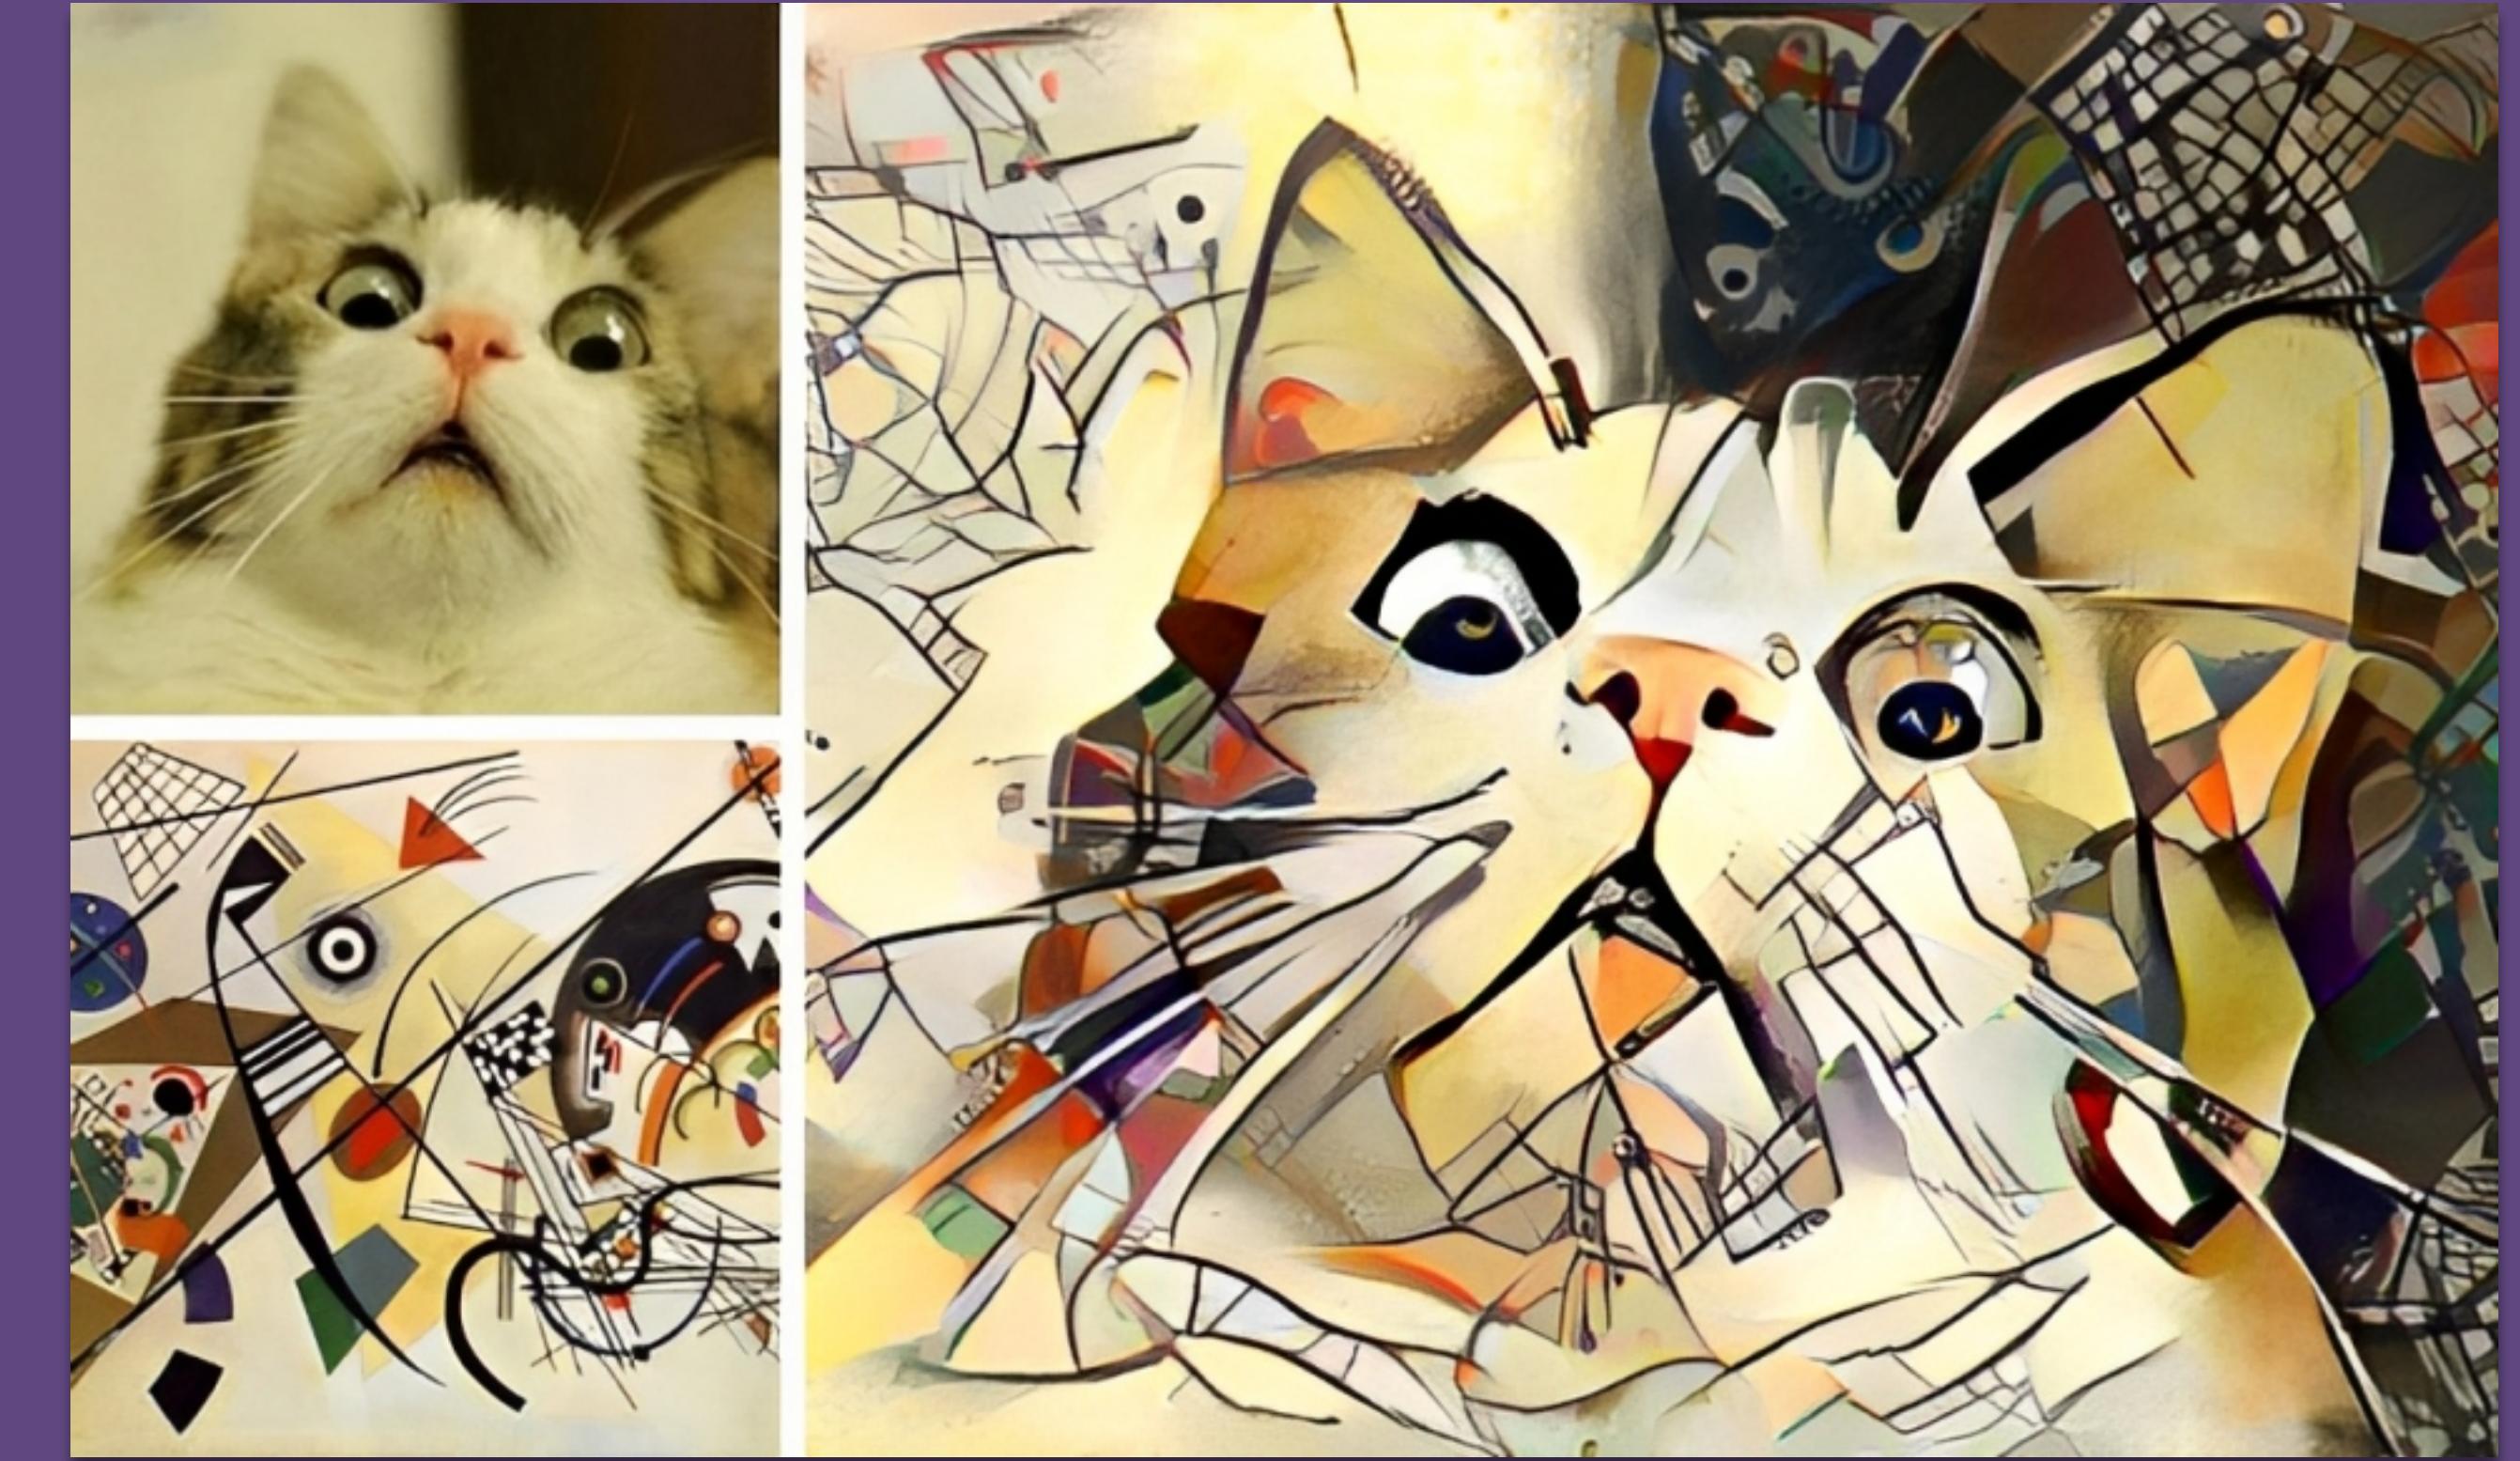


# Tutorial 3: Convolutional Neural Networks

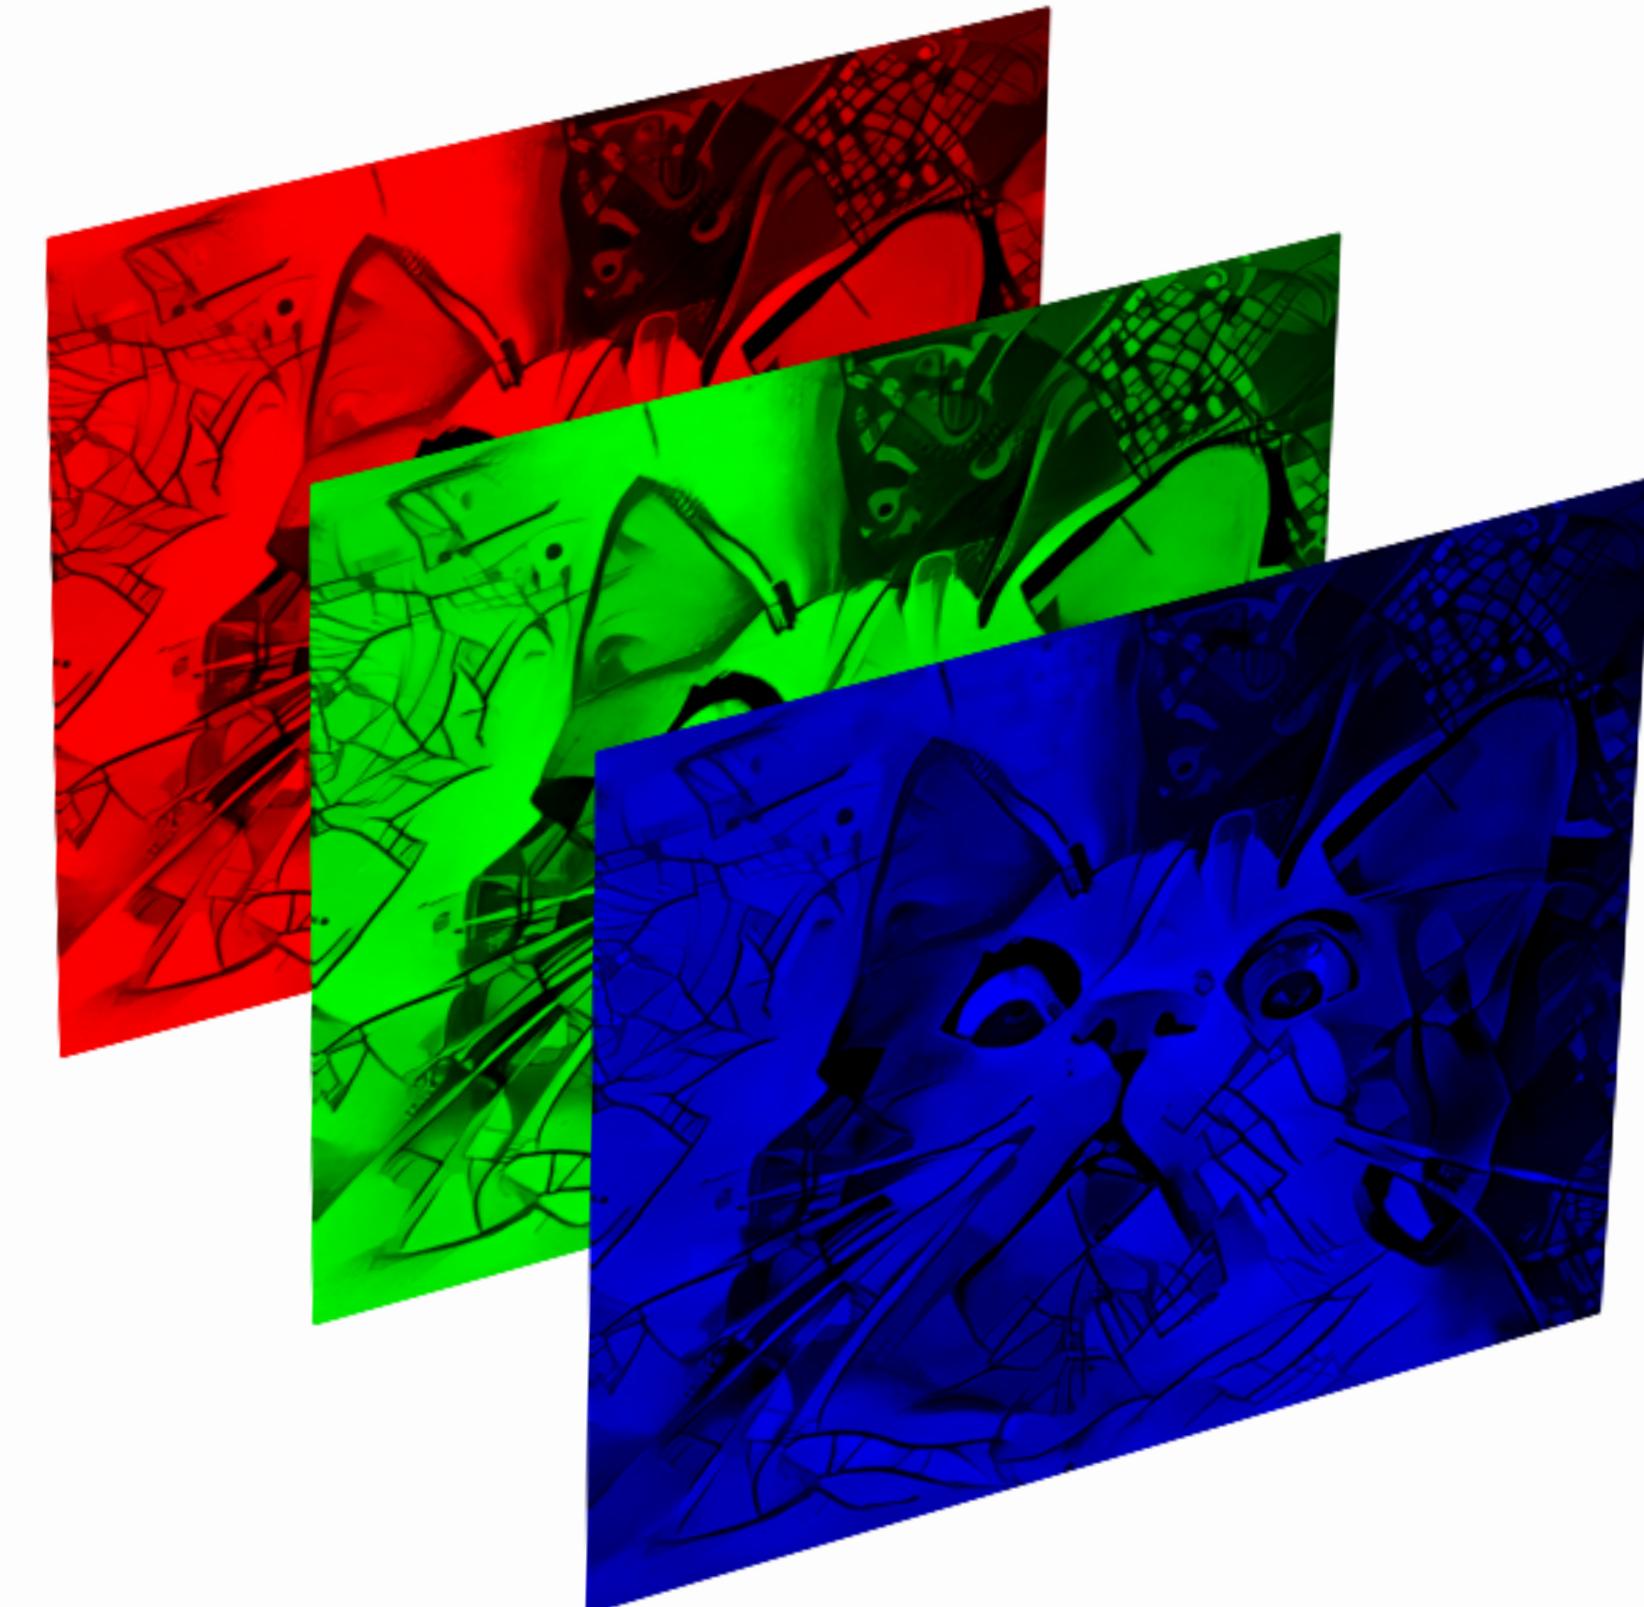
---



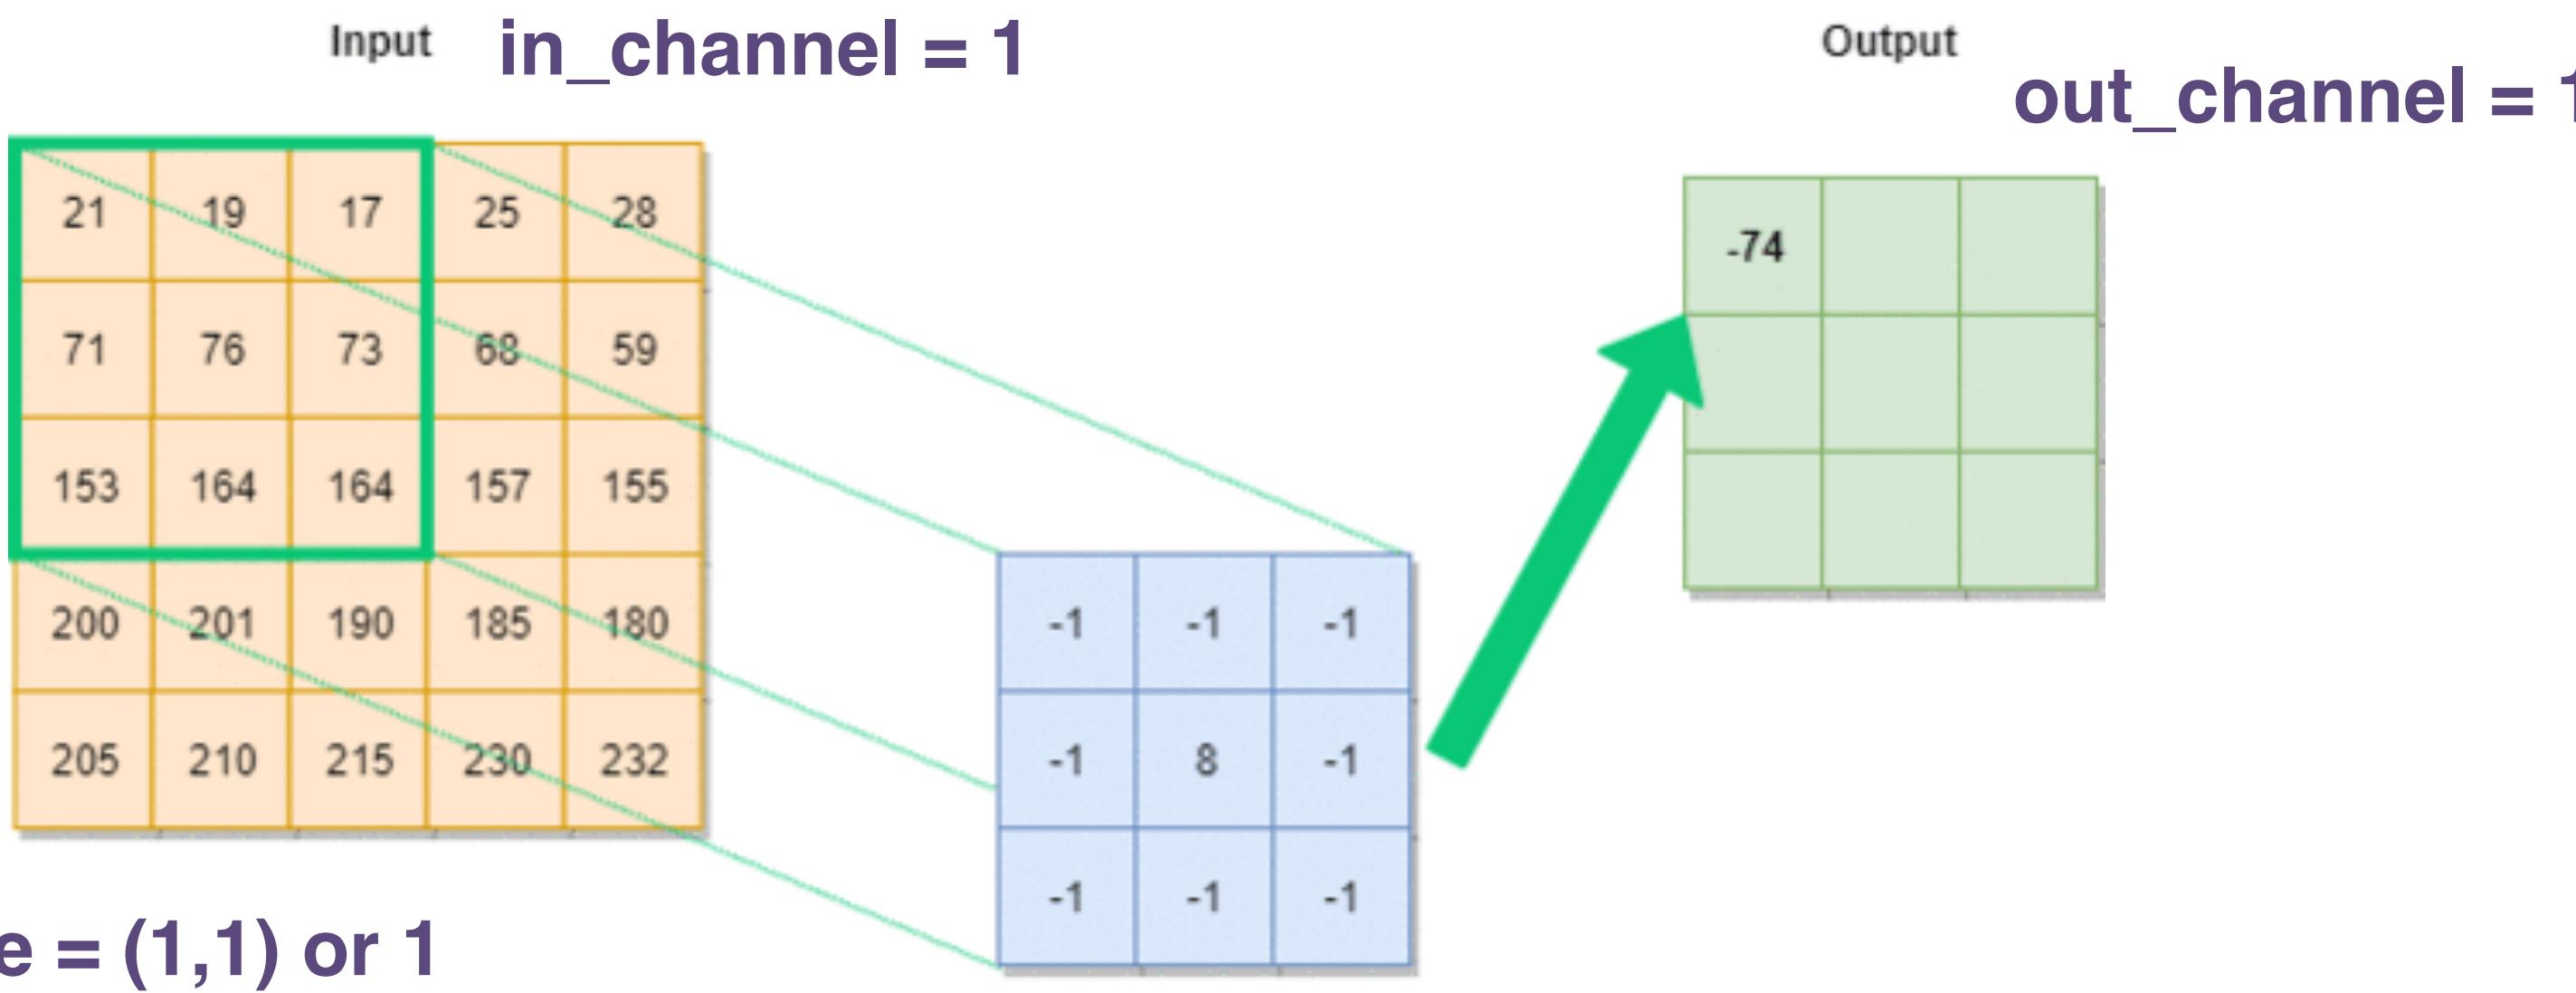
**Practical Deep Learning for Science**  
**10 April, 2025**

# CNN Vocabulary

# Channels



Not necessarily always 3

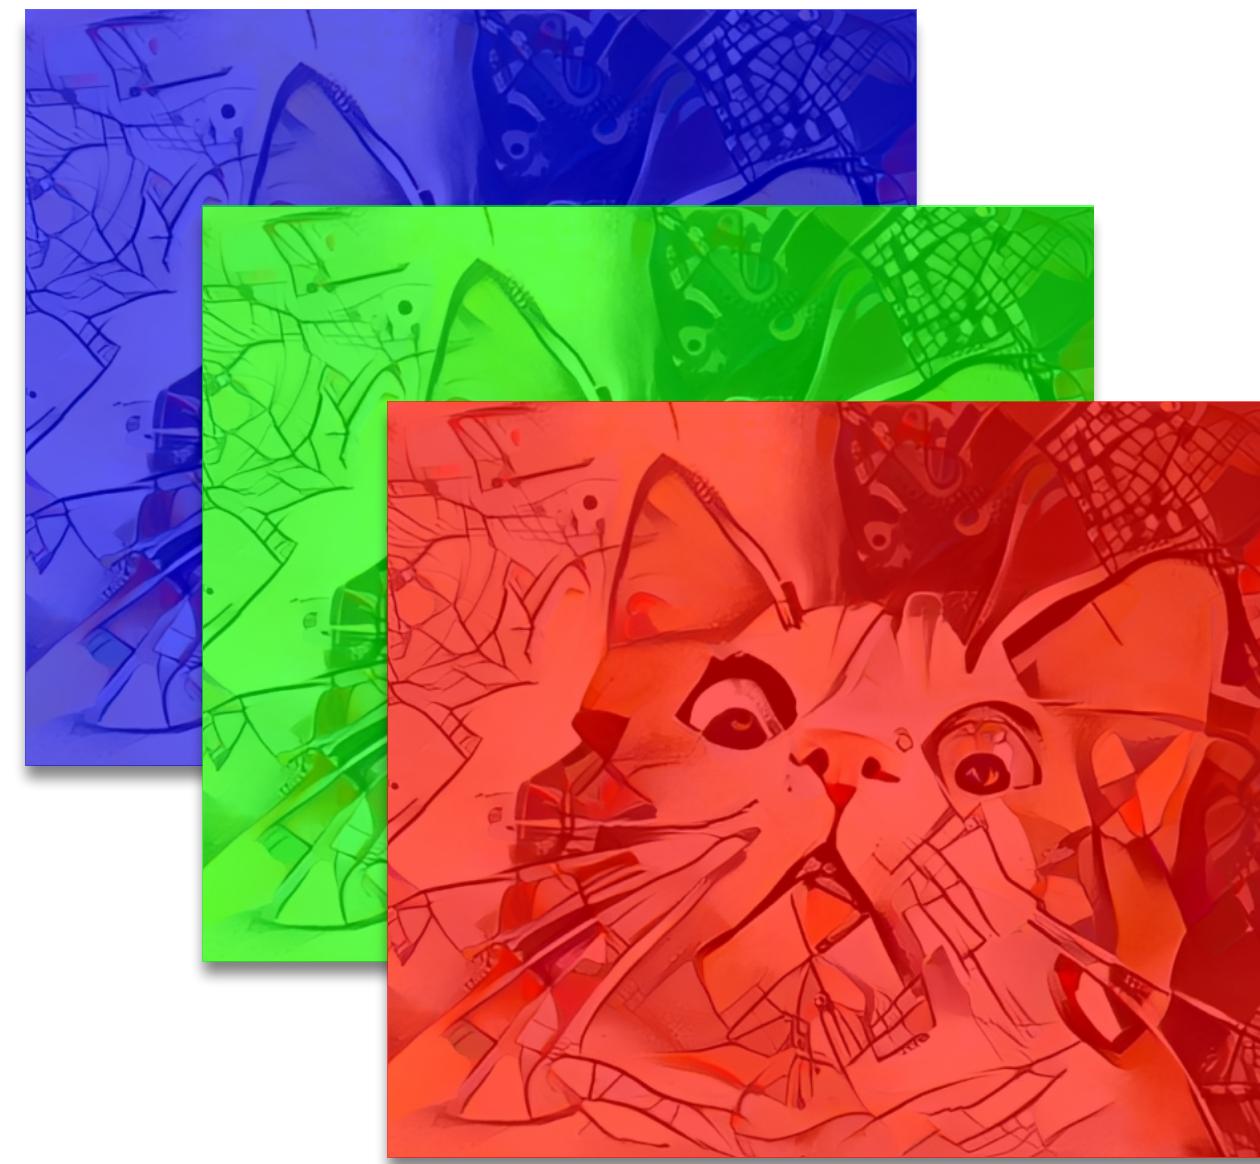


AIGeekProgrammer.com © 2019

**padding = ‘valid’ or 0**

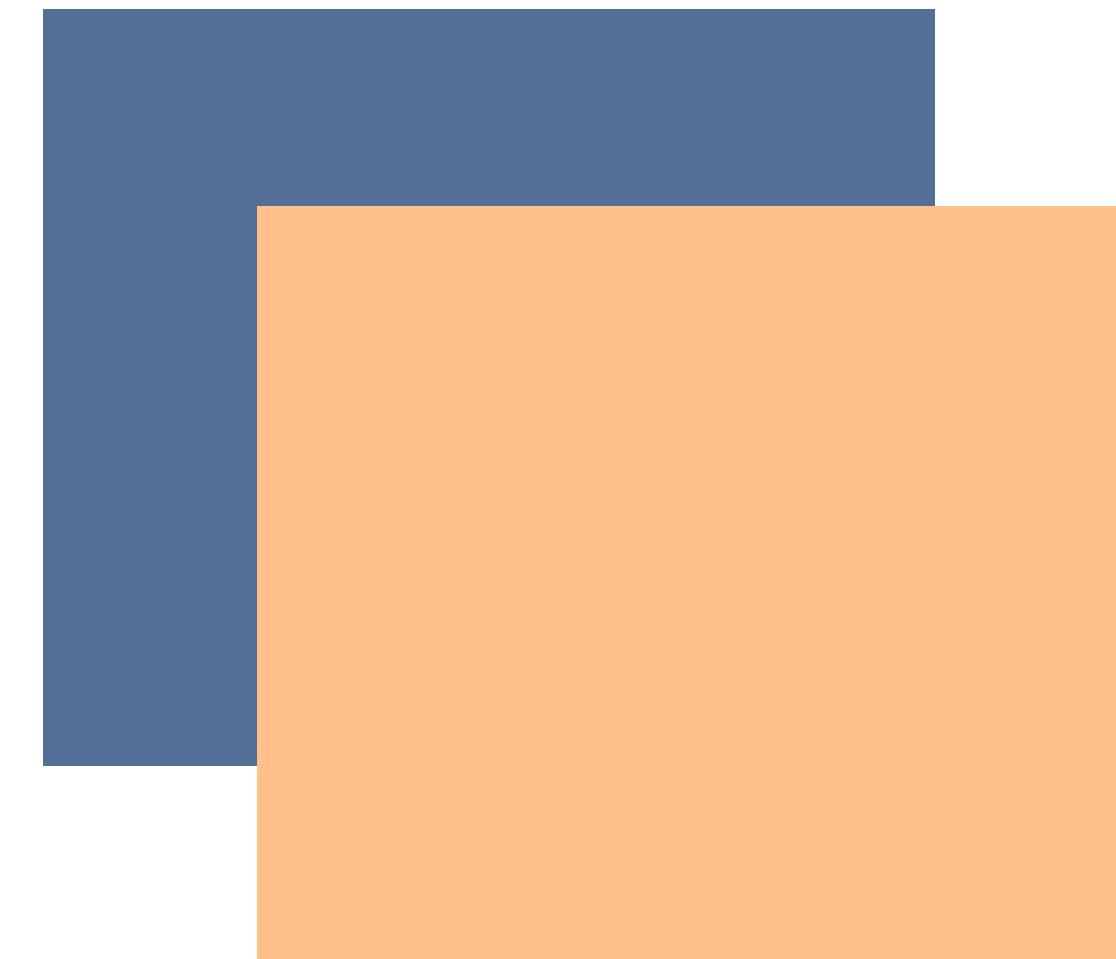
- **in\_channels**
- **out\_channels**
- **kernel\_size**
- **stride**
- **padding**

# Convolution with multiple channels



3 channels

Kernel = (5, 5) →

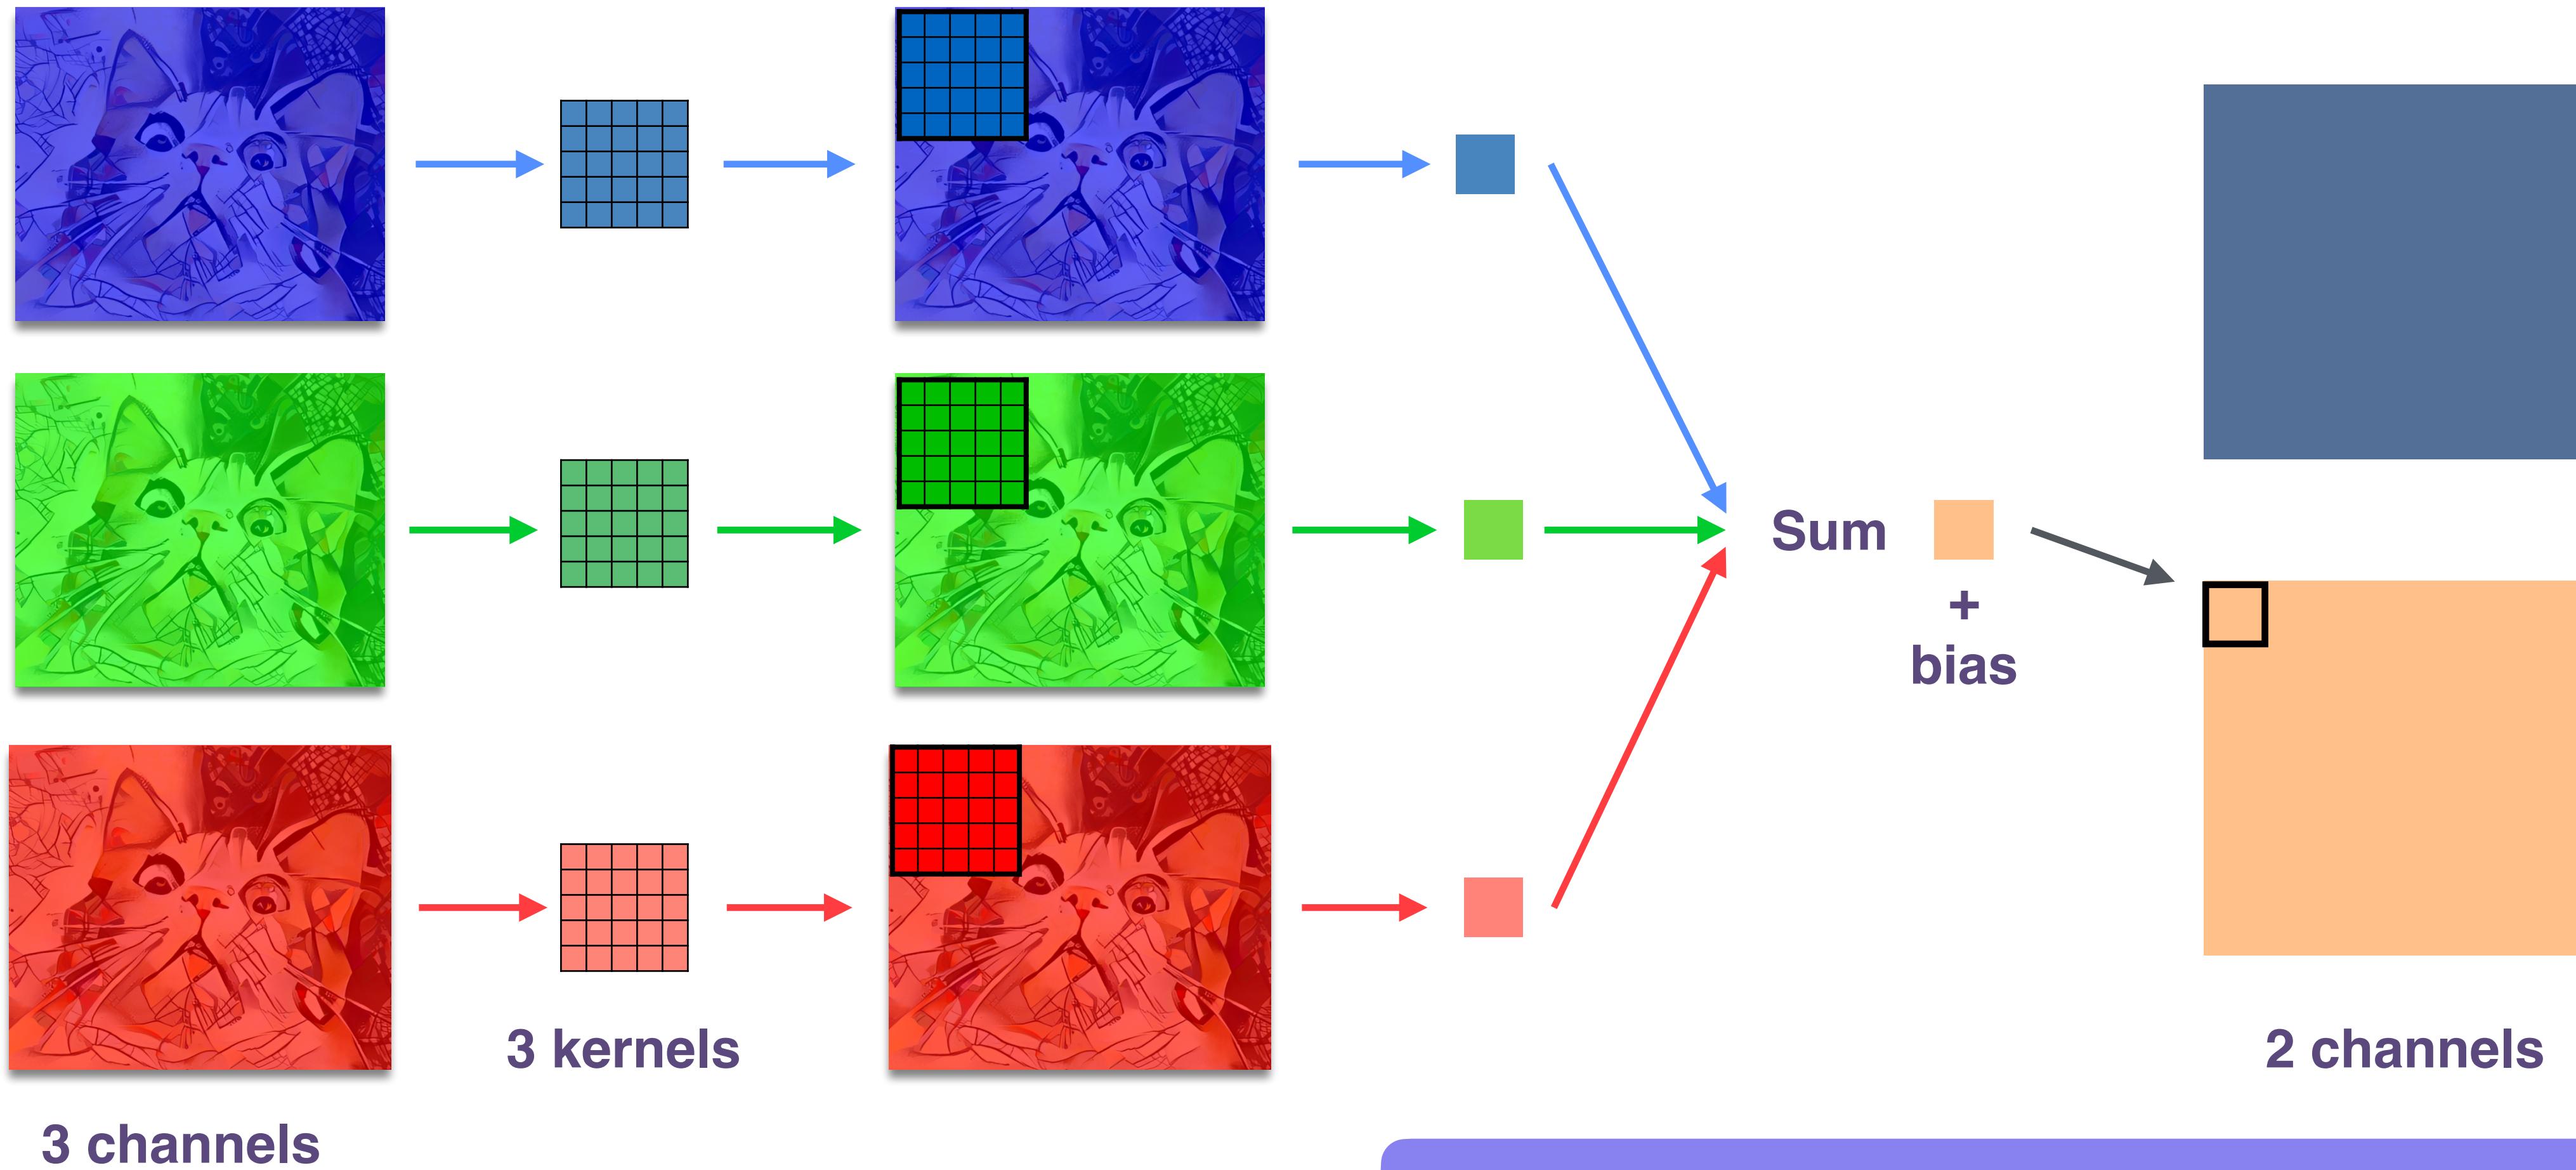


2 channels

*Ignore padding, stride etc.*

# Convolution with multiple channels

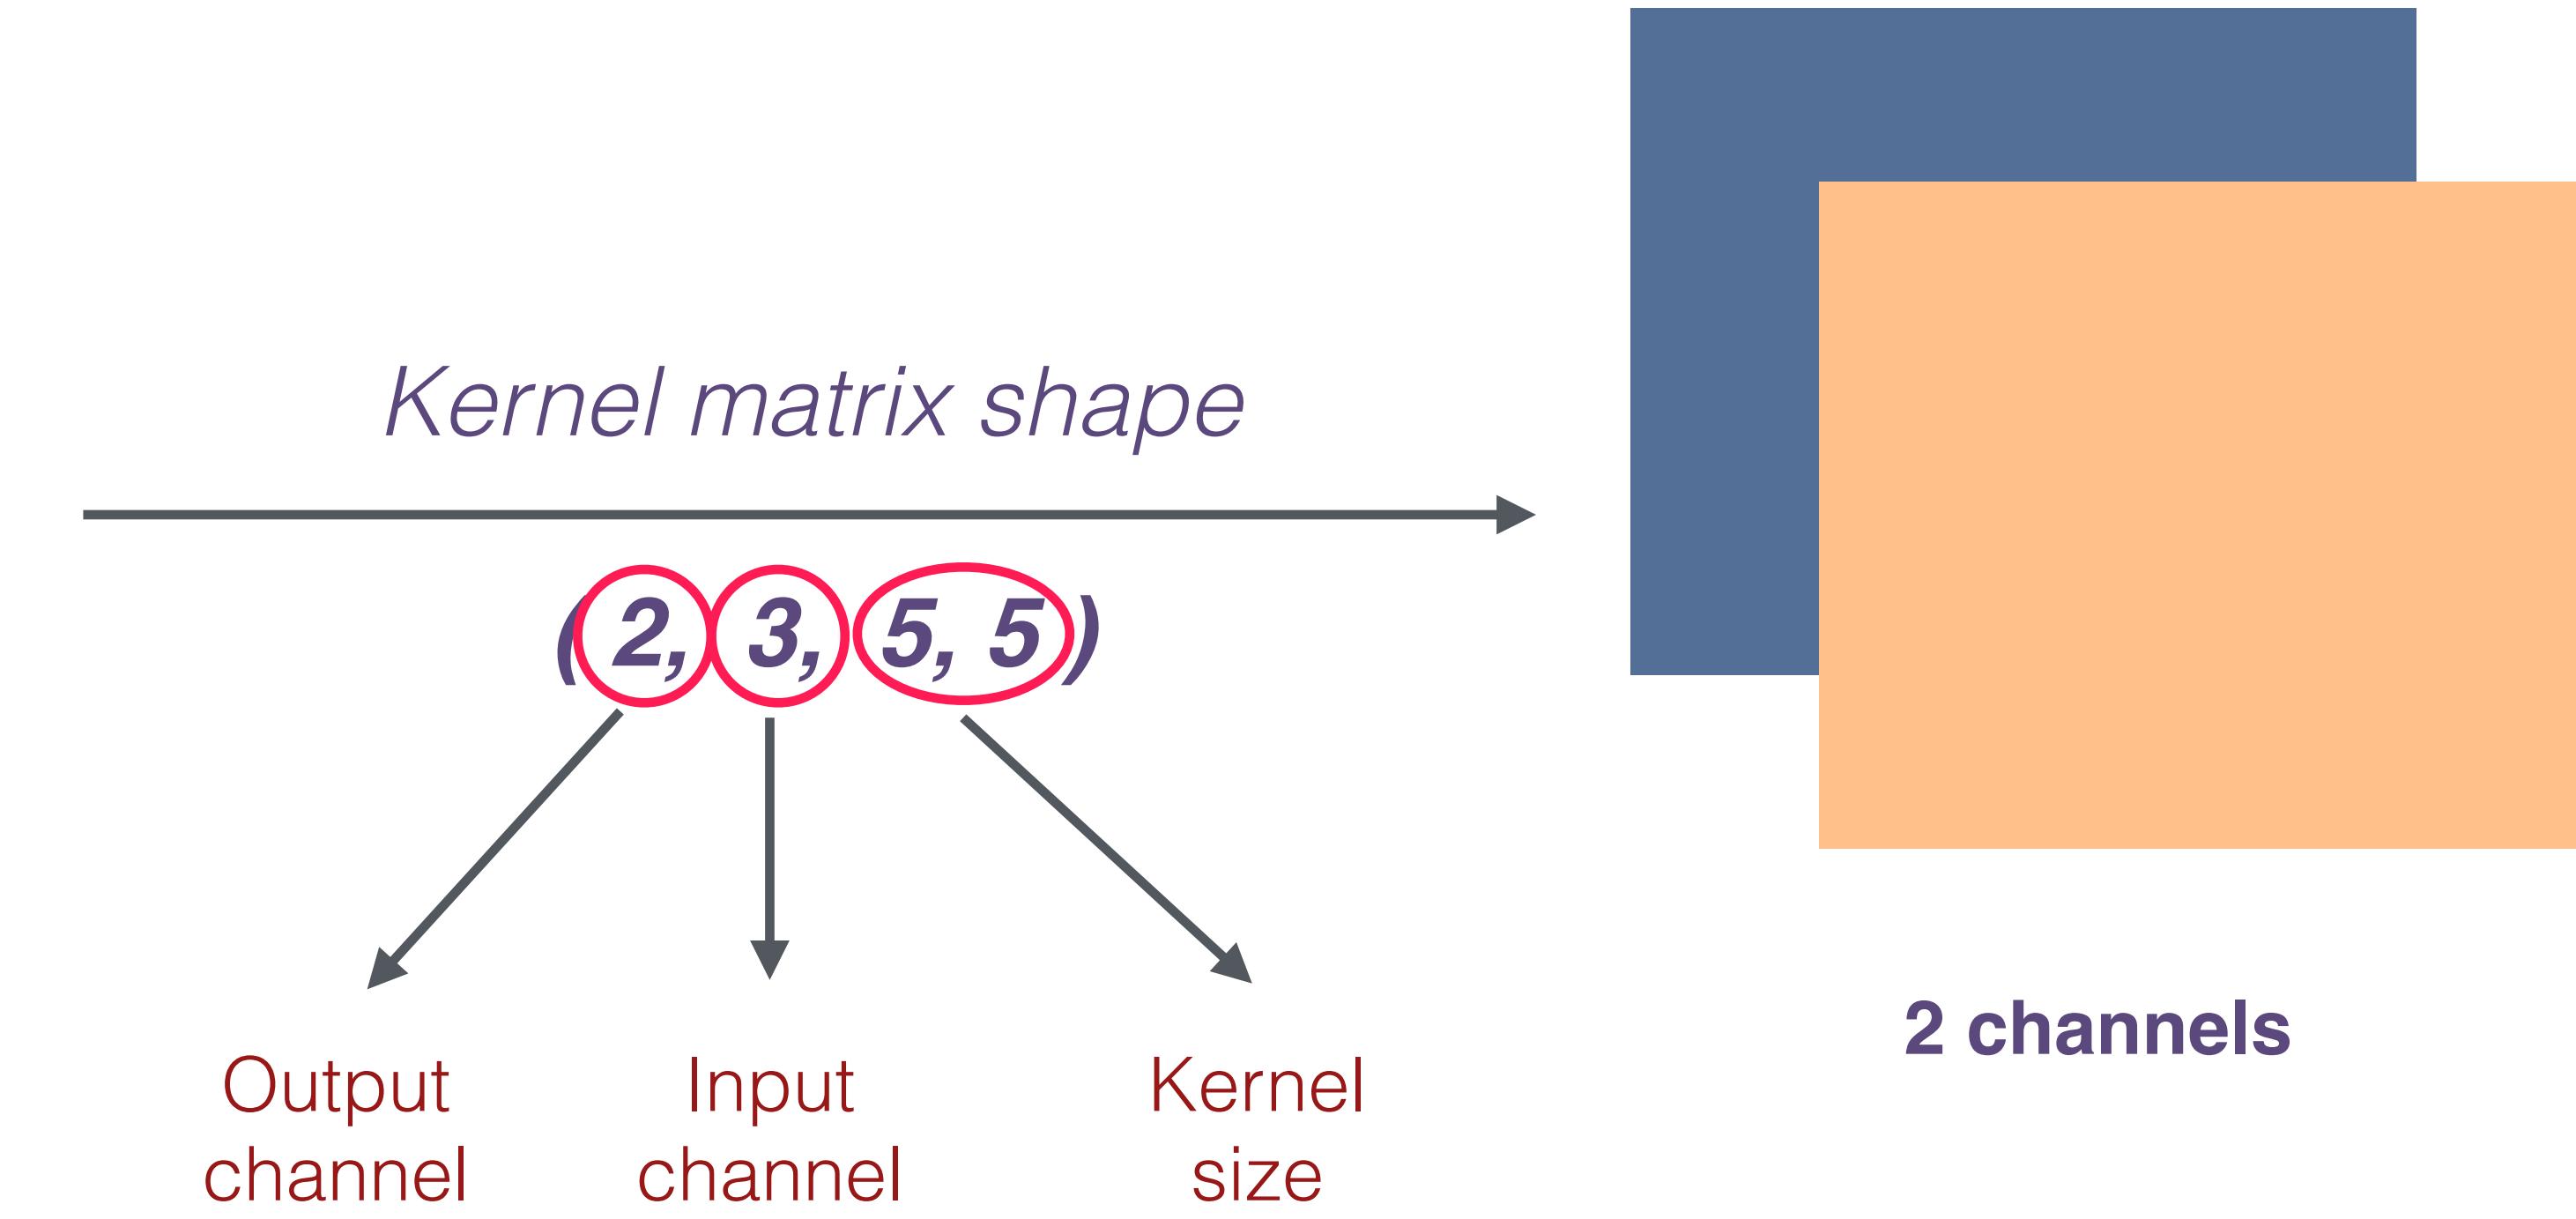
*Think in terms of one output pixel  
Cause we'll just move the kernel around for the full image*



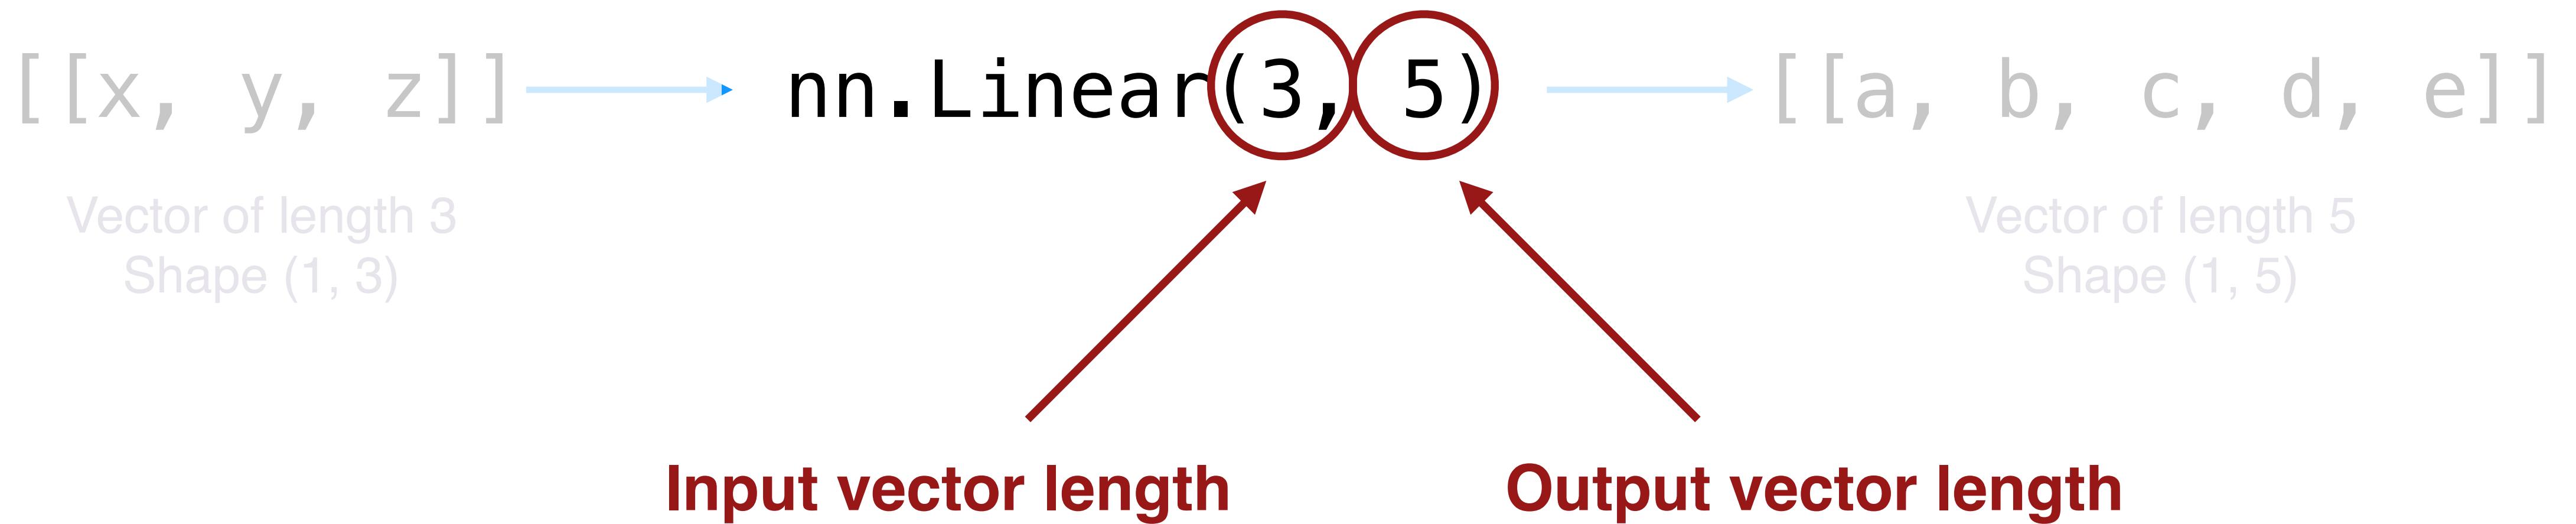
# Convolution with multiple channels



**3 channels**



# NHWC vs NCHW



$$(x, y, z) \begin{matrix} 1 \times 3 \\ \times 3 \end{matrix} \begin{pmatrix} w_{11} & w_{12} & w_{13} & w_{14} & w_{15} \\ w_{21} & w_{22} & w_{23} & w_{24} & w_{25} \\ w_{31} & w_{32} & w_{33} & w_{34} & w_{35} \end{pmatrix}_{3 \times 5} + \begin{pmatrix} b_1 & b_2 & b_3 & b_4 & b_5 \end{pmatrix}_{1 \times 5} = \begin{pmatrix} a & b & c & d & e \end{pmatrix}_{1 \times 5}$$

A is a `torch.Tensor()`. Previously we saw a 2D tensor

A.shape			
1D tensor	(2, )	→	<code>nn.Linear(?, 5)</code>
2D tensor	(1, 3)	→	<code>nn.Linear(3, 5)</code>
3D tensor	(2, 3, 4)	→	<code>nn.Linear(?, 5)</code>
4D tensor	(2, 3, 4, 5)	→	<code>nn.Linear(?, 5)</code>

A is a `torch.Tensor()`. Previously we saw a 2D tensor

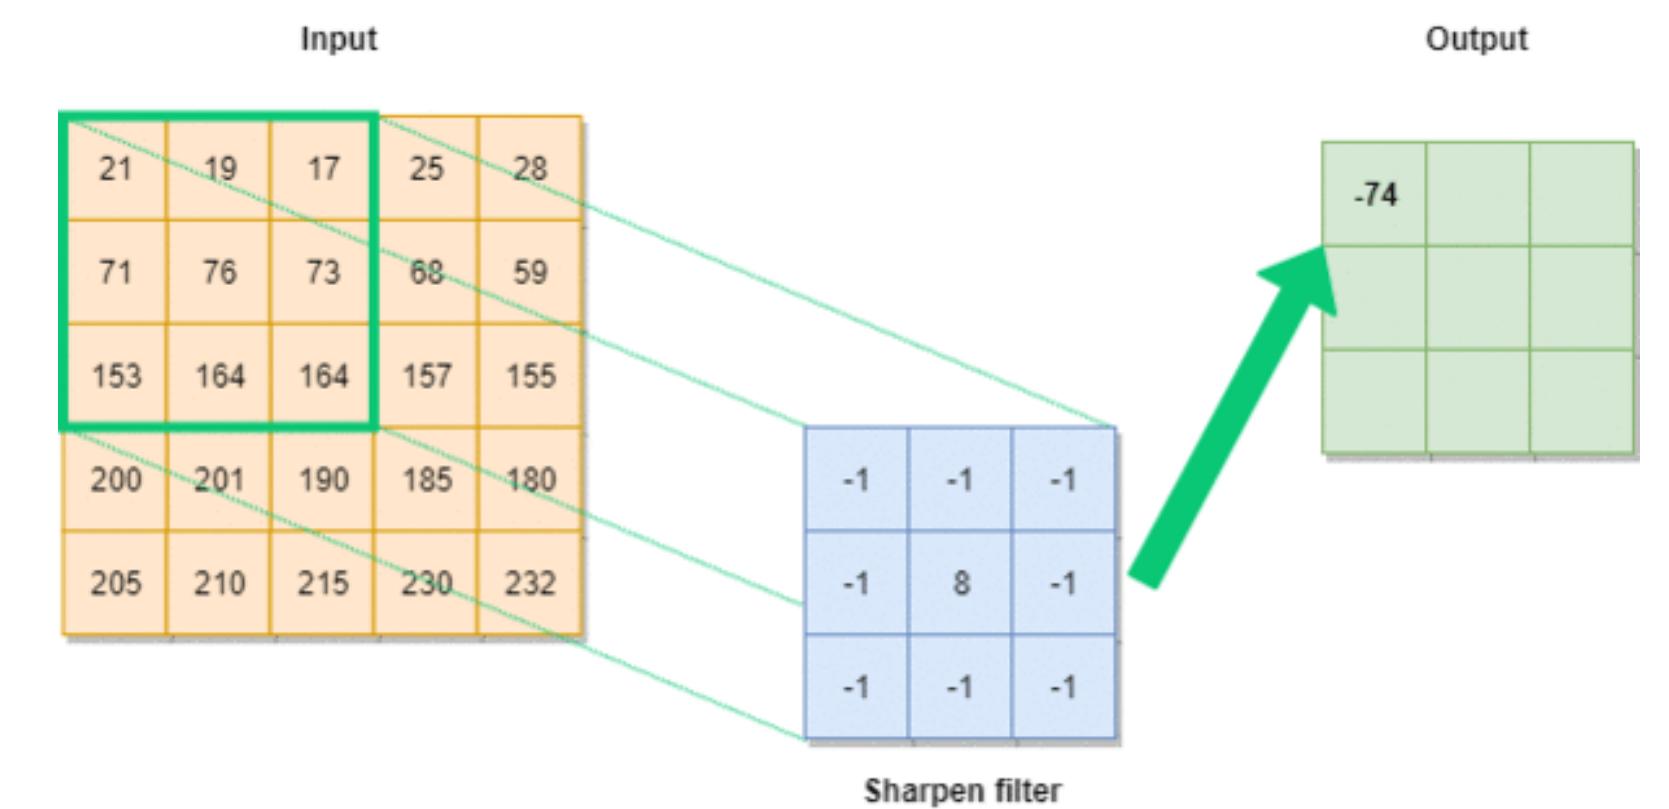
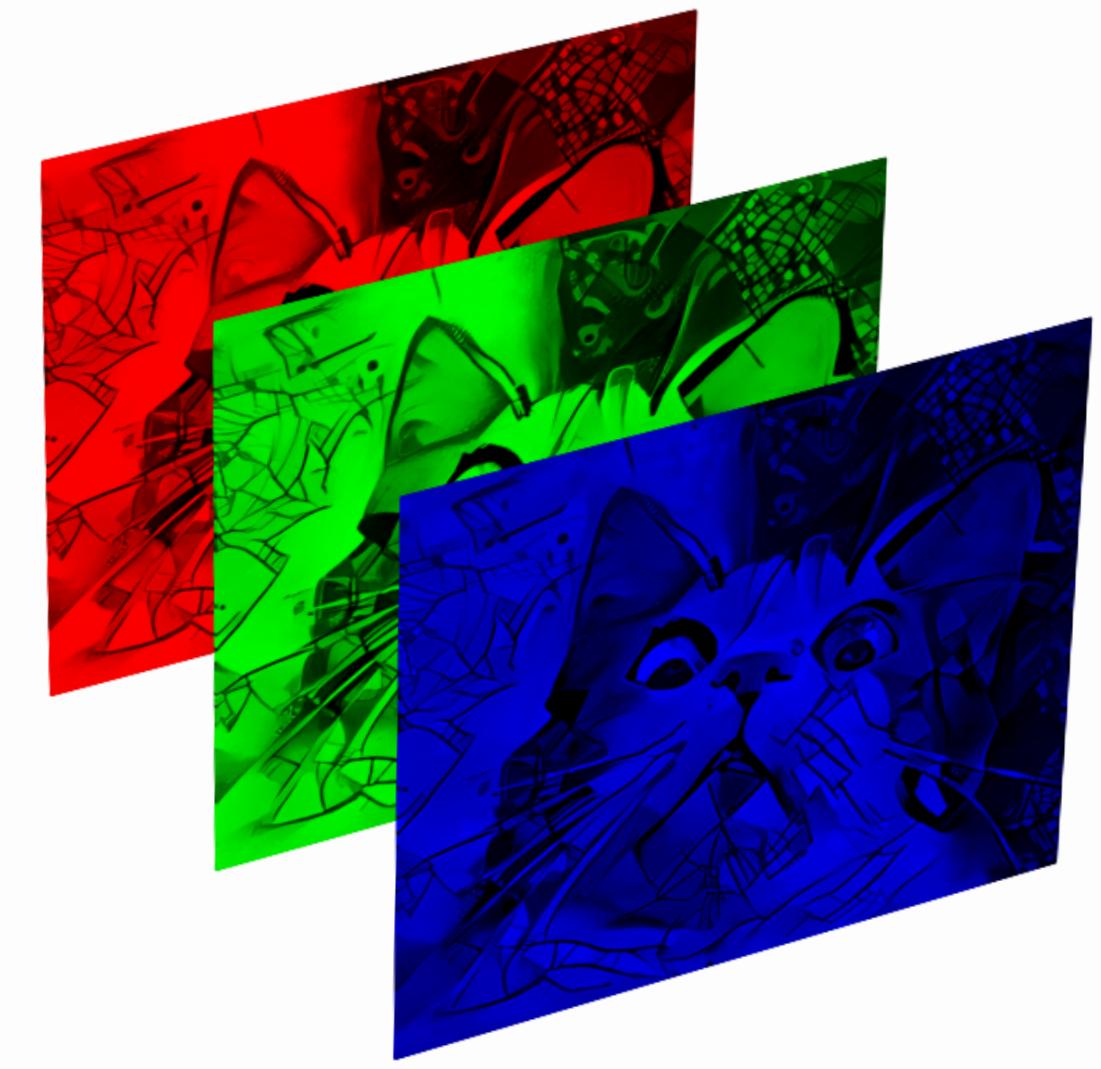
In torch, It'll always act on the last dimension

## A.shape

1D tensor	(2,)	→	<code>nn.Linear(2, 5)</code>
2D tensor	(1,3)	→	<code>nn.Linear(3, 5)</code>
3D tensor	(2,3,4)	→	<code>nn.Linear(4, 5)</code>
4D tensor	(2,3,4,5)	→	<code>nn.Linear(5, 5)</code>

# What about Conv2D

- ♦ In general, images will have shape
  - **(batch\_size, H, W, 3)** (3 channels - RGB)
- ♦ Kernel will have shape, say
  - **(Kernel\_H, Kernel\_W)**
- ♦ With this setup, we will act the kernel on **(W, 3)**, but it should act on **(H, W)**
- ♦ So, we need to change the image format a bit



# NCHW vs NHWC

- **N** - number of example in the batch
- **C** - channel number
- **H** - height
- **W** - width
- Torch expects data to be in **(N,C,H,W)** format

**Permute and reshape are two completely different operations!**

- An **RGB** image of size **100x50**

What we want

```
>>> img_tensor.shape  
torch.Size([1, 3, 100, 50])
```

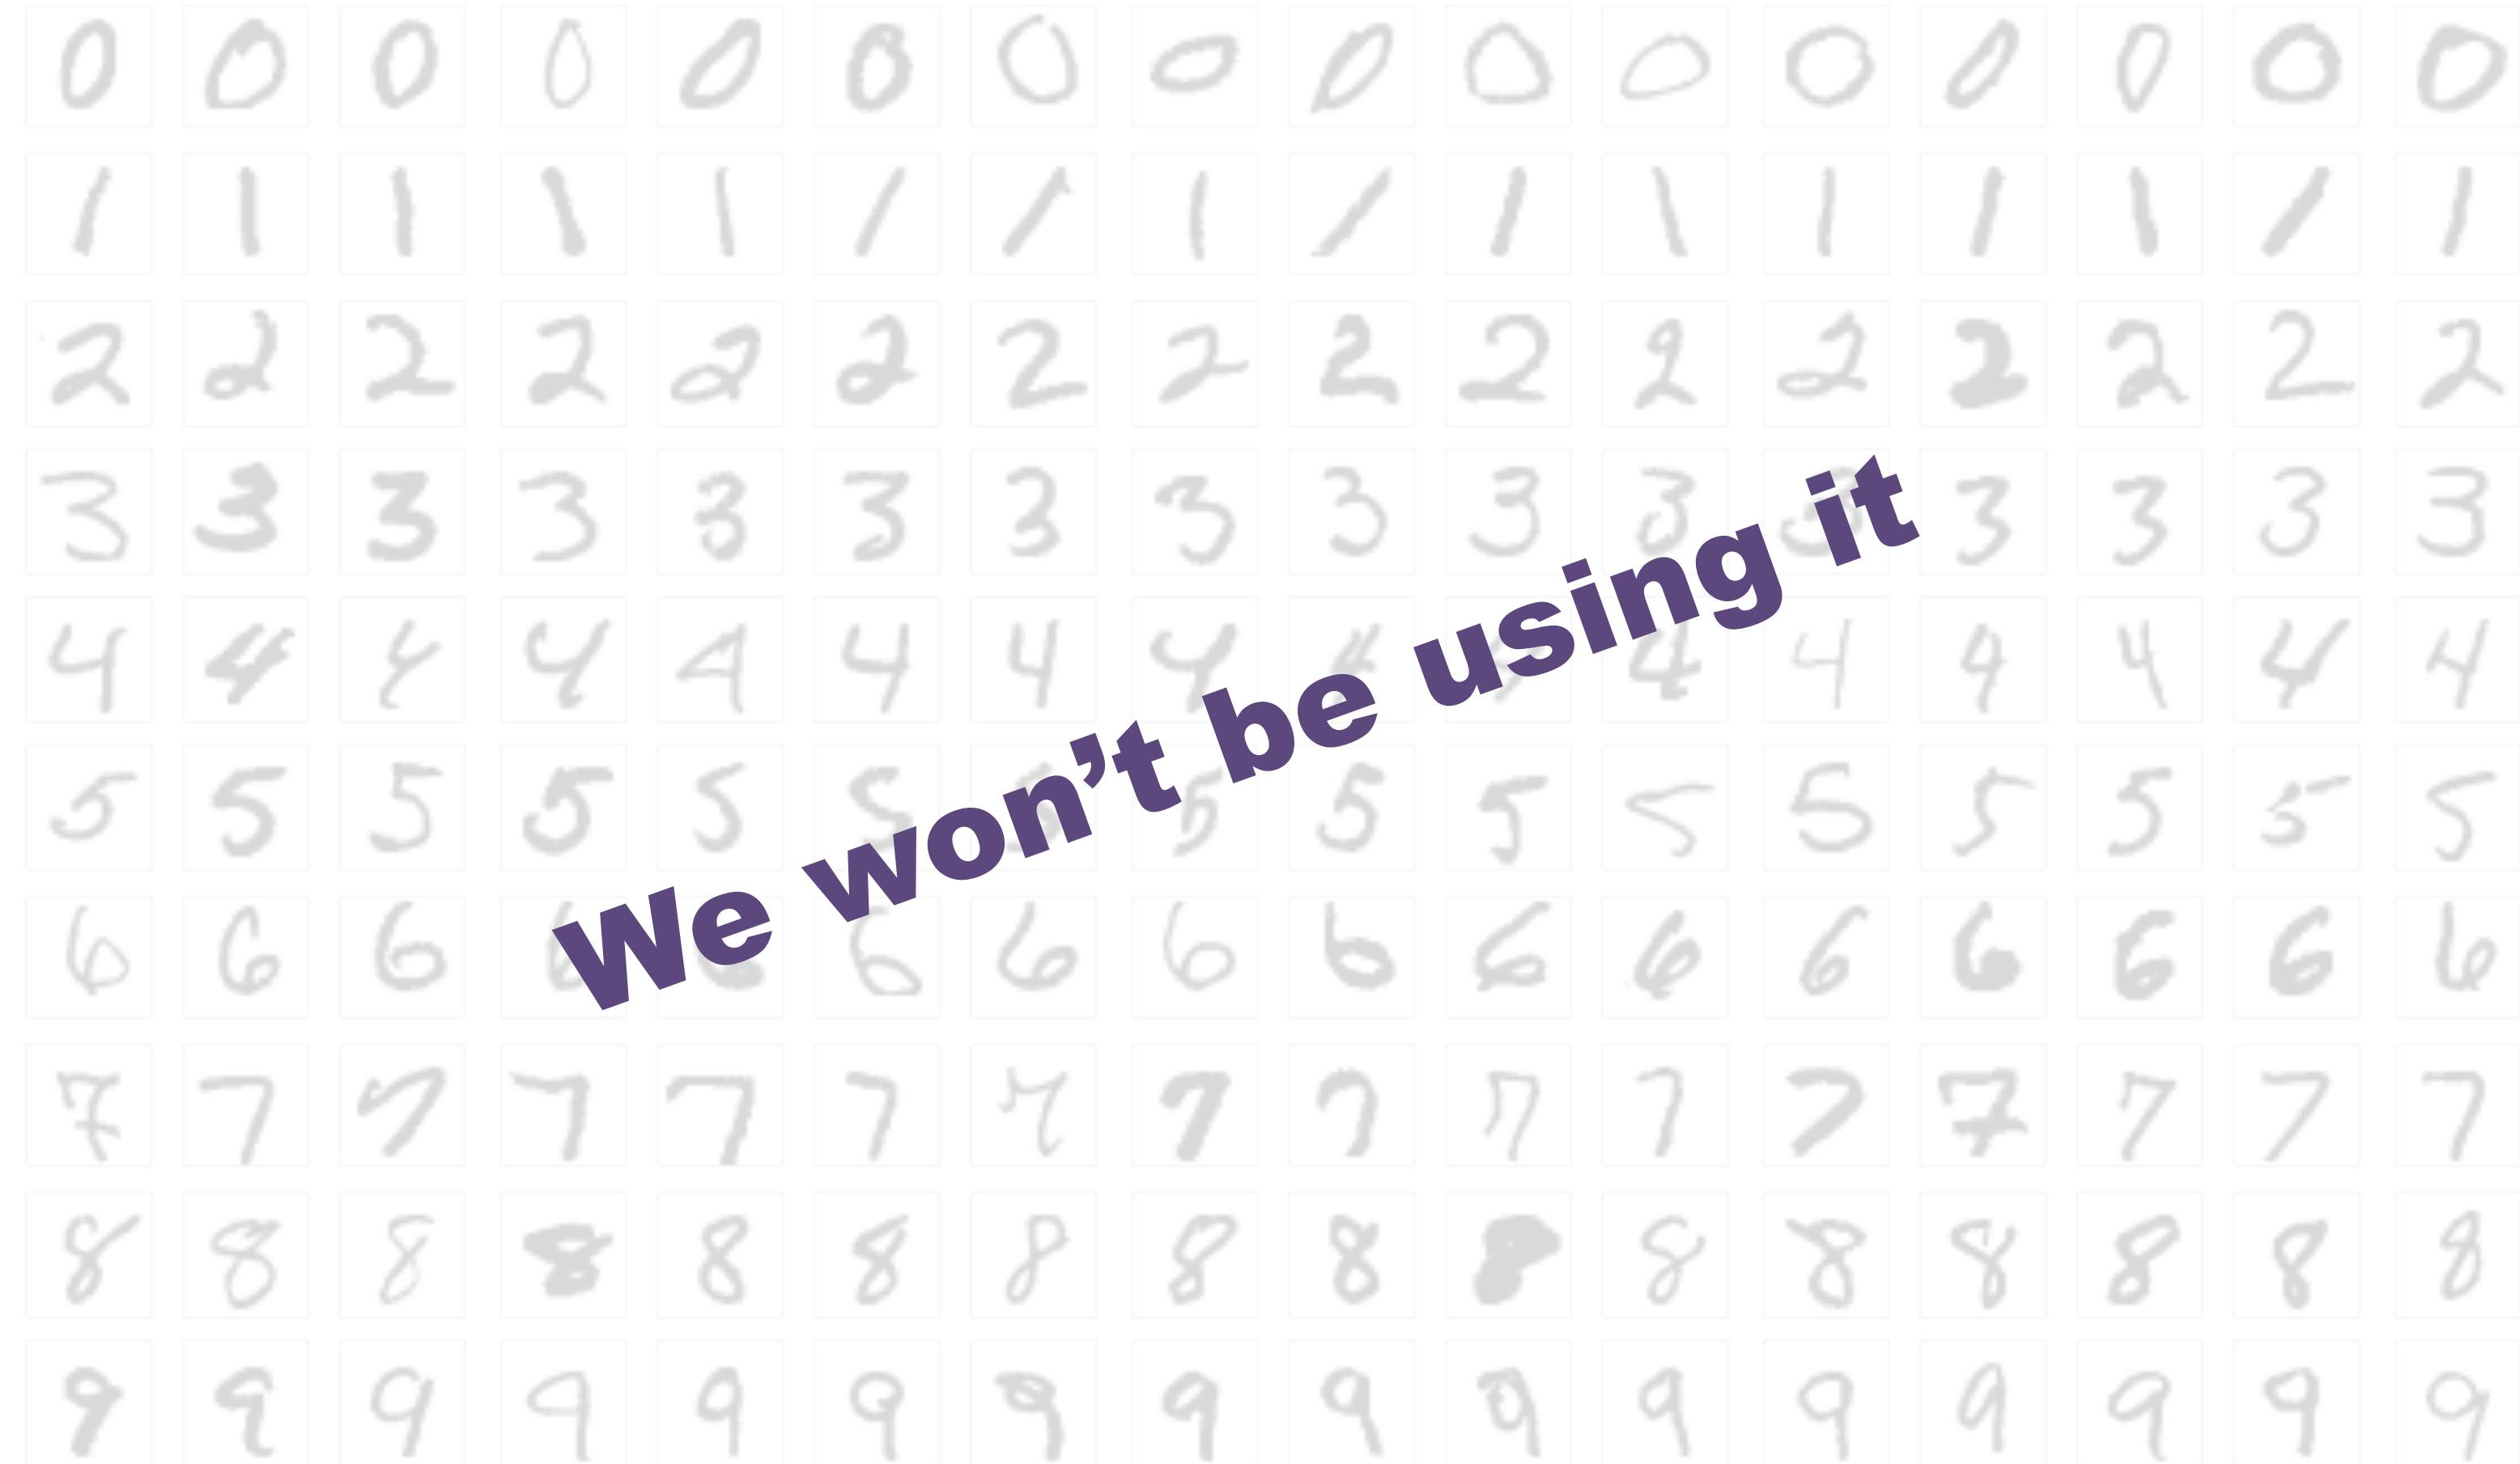
What we do not want

```
>>> img_tensor.shape  
torch.Size([1, 100, 50, 3])
```

Fix

```
>>> img_tensor = img_tensor.permute(0,3,1,2)  
>>> img_tensor.shape  
torch.Size([1, 3, 100, 50])
```

# Today's dataset



## MNIST

- Handwritten digits
- 28x28
- Grayscale (1 channel)
- Pretty old (1994)
- Quite simple

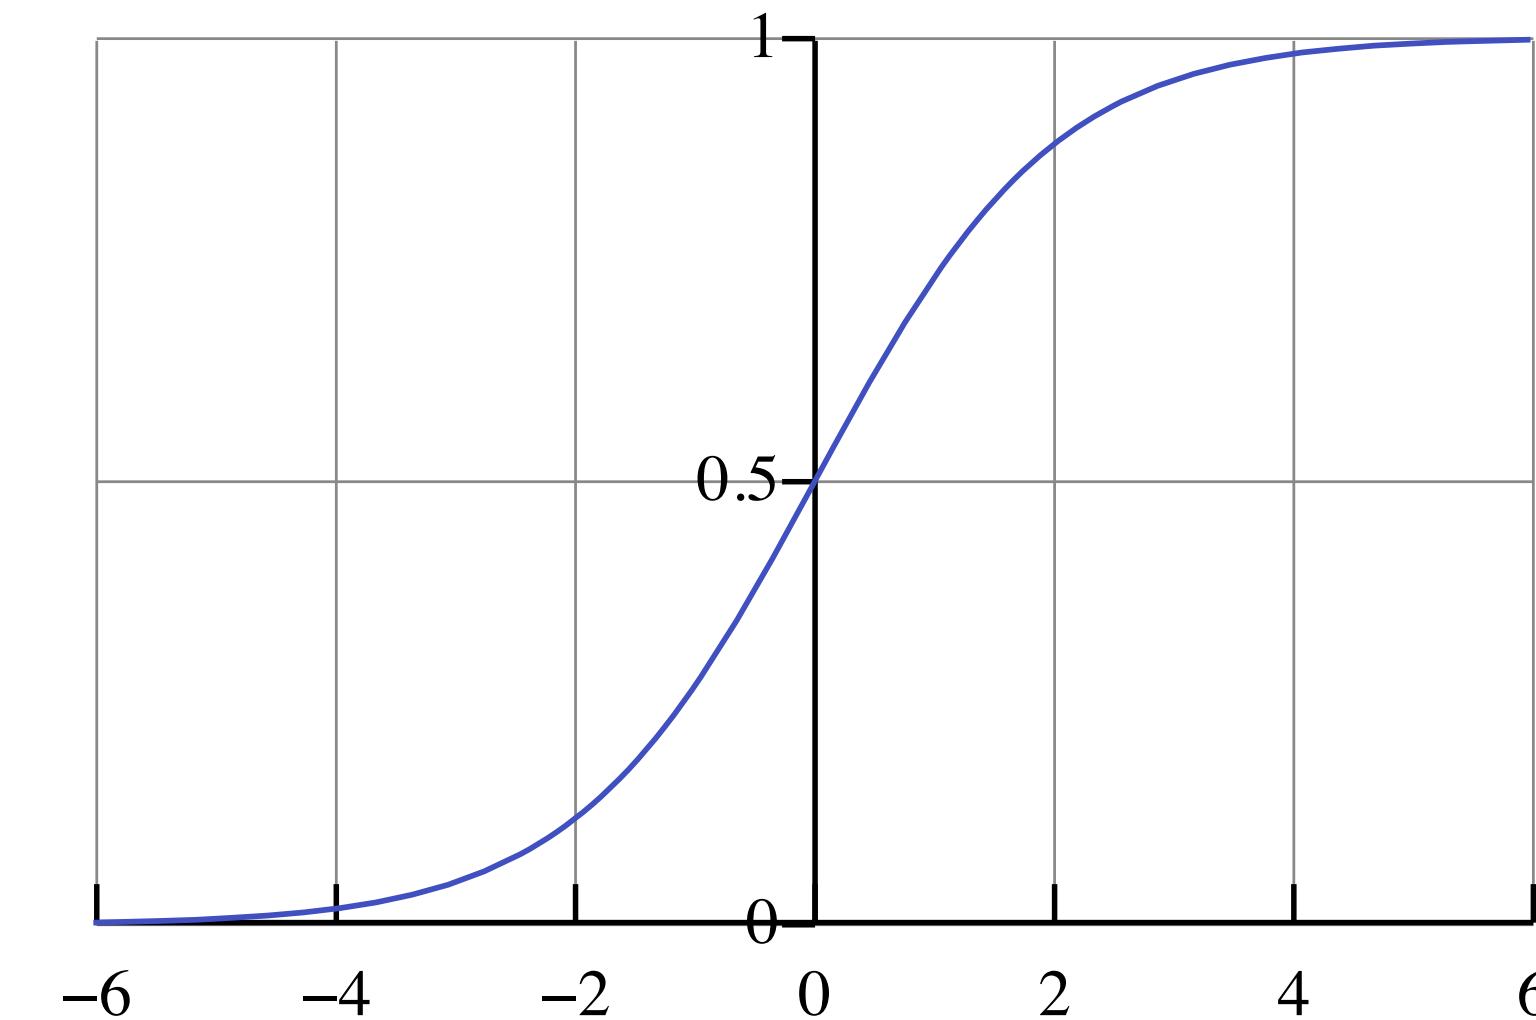
# Fashion-MNIST

- 28 x 28
- grayscale images (1 channel)
- 10 classes
- Slightly more complex than MNIST

# Activations and losses

# Binary classifications

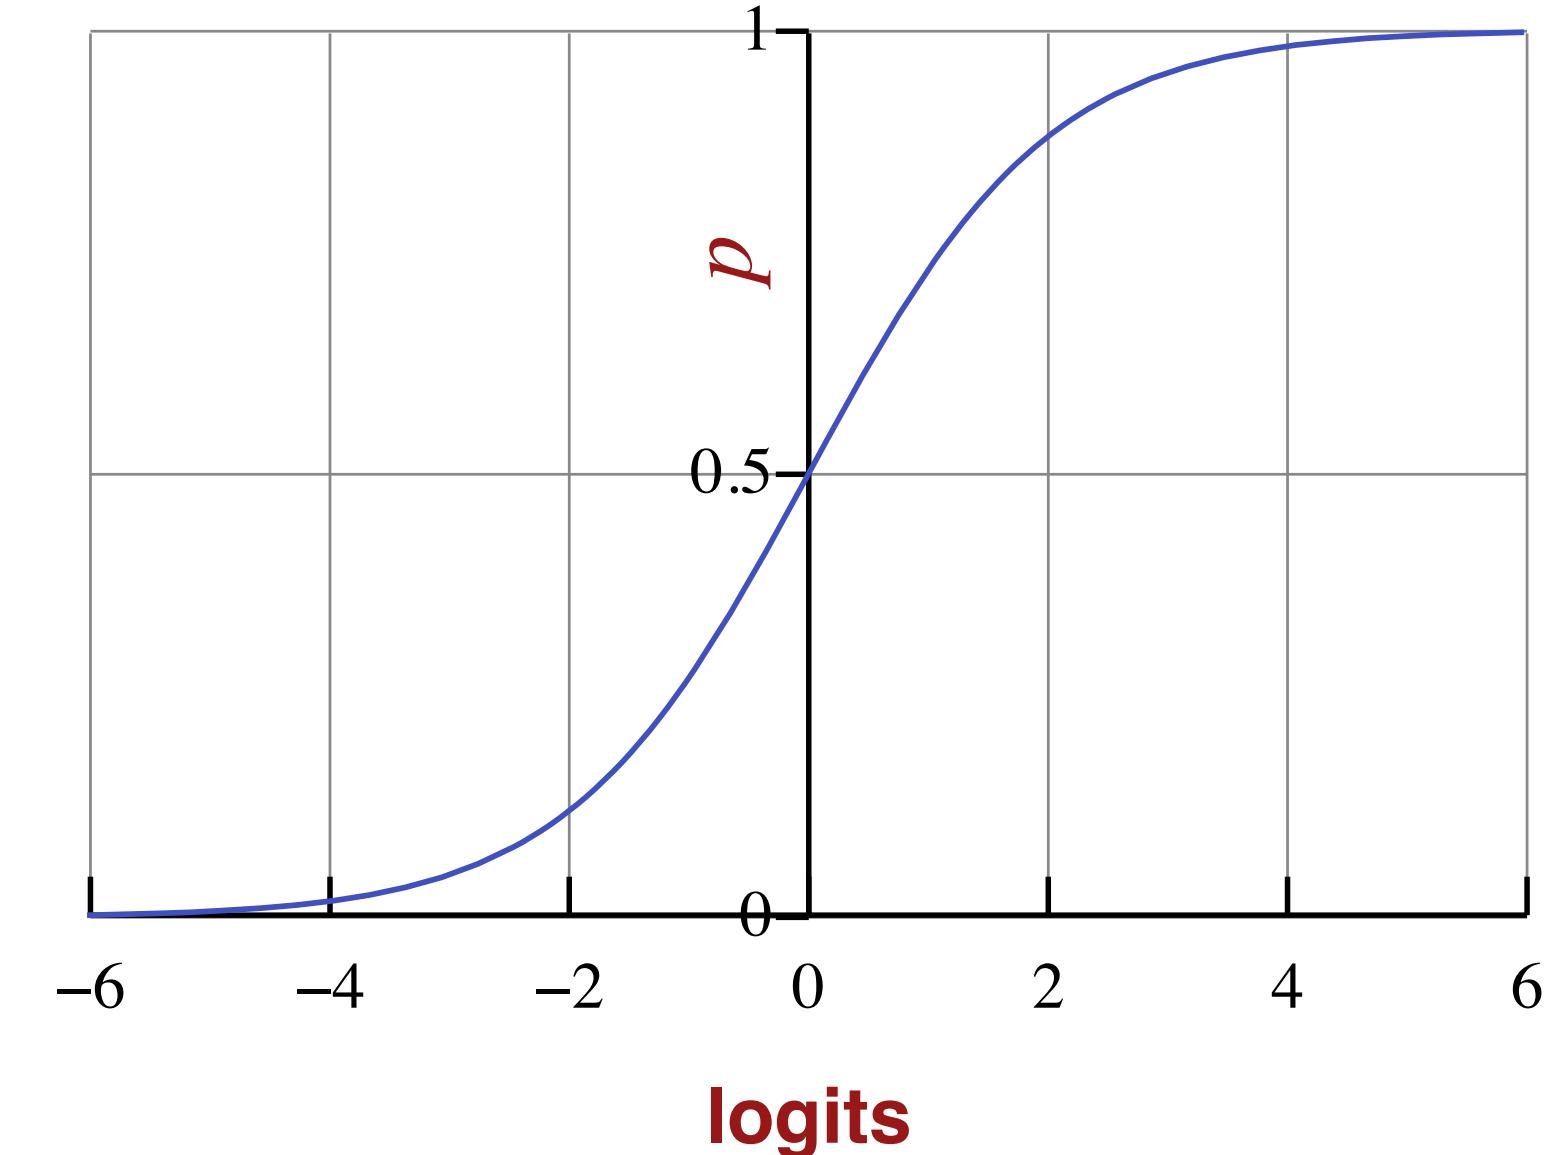
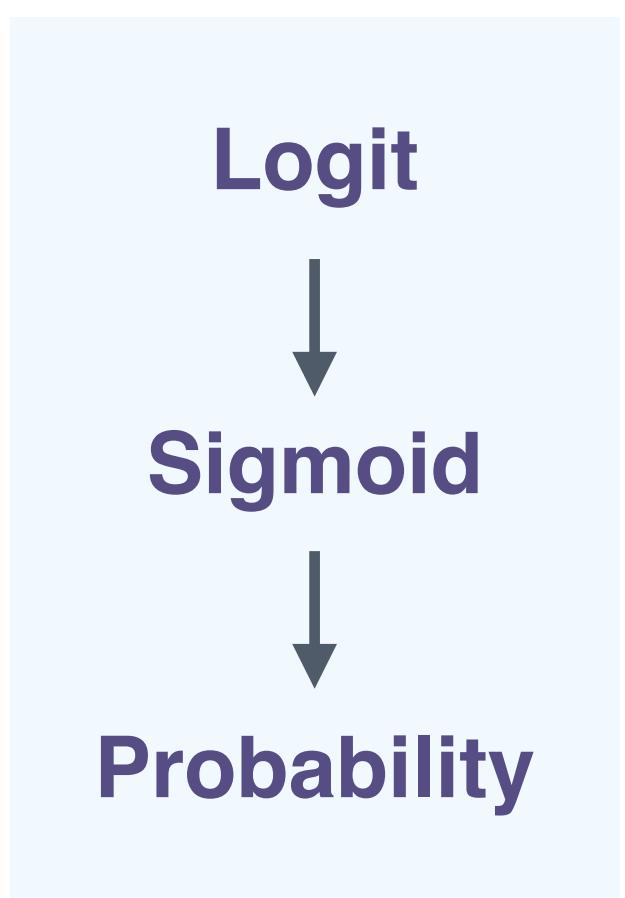
- **Class 0 vs Class 1**
  - NN Prediction = **probability of class 0** ( $p$ )
  - $\text{Probability of class 1} = 1 - p$ 
    - No need to predict this
  - Activation: `torch.nn.Sigmoid()`



# Binary classifications: Loss

- **What we want -**

- End the network with `Linear() → Sigmoid()`
- Compute the loss using -
  - `torch.nn.BCELoss(pred_prob, target)`



- **More common practice -**

- End the network with `Linear()`
- Compute the loss using -
  - `torch.nn.BCEWithLogitLoss(pred_logit, target)`
  - Internally applies `Sigmoid()`
  - Numerically more stable

- **During inference -**

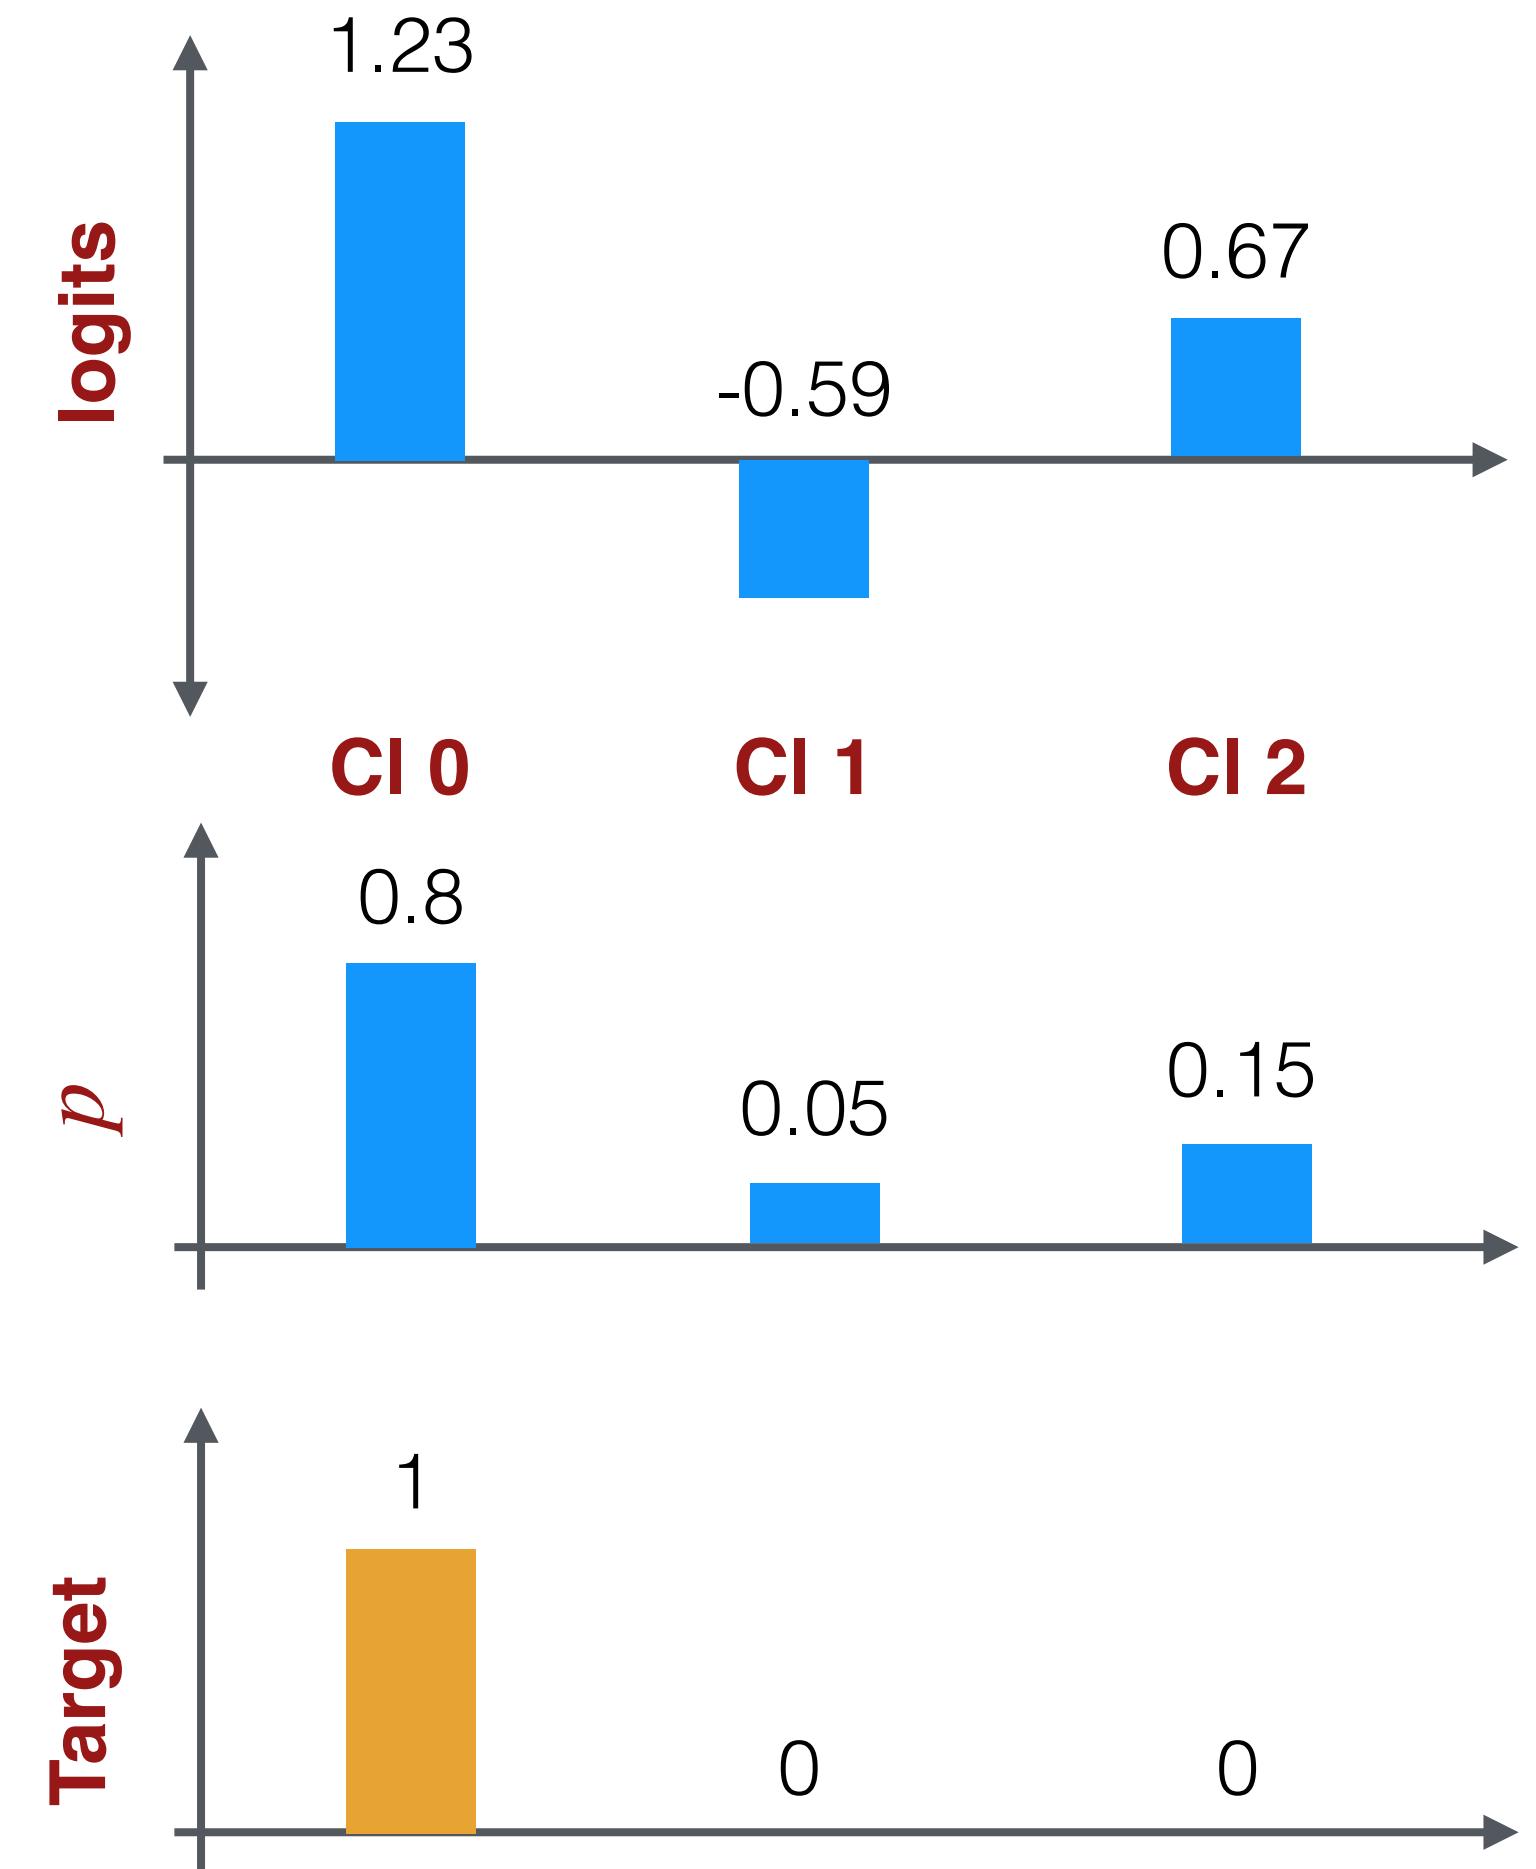
- $p > 0.5 \implies$  class 0
- $logit > 0 \implies$  class 0

# Multi-class classifications

- **Option one -**

- End the network with `Linear() → Softmax()`
- Compute the loss using -
  - `torch.nn.NLLLoss(pred_prob, target)`

$$p_i = \frac{e^{x_i}}{\sum_i e^{x_i}}$$



- **More common practice -**

- End the network with `Linear()`
- Compute the loss using -
  - `torch.nn.CrossEntropy(pred_logit, target)`
  - Internally applies `Softmax()`
  - Numerically more stable

- **During inference -**

- Class = `argmax(prob)`
- Class = `argmax(logit)`